

DeepLense: Strong Gravitational Lensing Classification

Google Summer of Code (GSoC) 2026 – Common Task Report

Author: Puttala Jeevan Kumar

Date: January 29, 2026

Organization: ML4Sci Foundation

Repository: [<https://github.com/Jeevang1-epic/ML4Sci-DeepLense-GSoC-Jeevan.git>]

1. Abstract

This report details the development of a deep learning solution for the ML4Sci "DeepLense" Common Task (Model I). The objective was to perform binary classification on simulated strong gravitational lensing images to distinguish between smooth dark matter halos (no_sub) and halos containing substructure (sub). Using a modified **ResNet-18** architecture with **Mixed Precision (AMP)** training, the model achieved a verified **ROC AUC score of 1.0000** on the test dataset. This report outlines the methodology, architectural decisions, and performance metrics validation.

2. Problem Statement

Dark matter substructure detection is a critical challenge in cosmology. Strong gravitational lensing offers a probe into these substructures. The task involves automating the classification of single-channel (grayscale) lensing images into two categories:

- **Class 0 (no_sub):** Lensing caused by a smooth dark matter halo.
- **Class 1 (sub):** Lensing significantly perturbed by dark matter sub-halos.

The primary evaluation metric for this task is the **Area Under the Receiver Operating Characteristic Curve (ROC AUC)**.

3. Methodology

3.1 Data Preprocessing & Augmentation

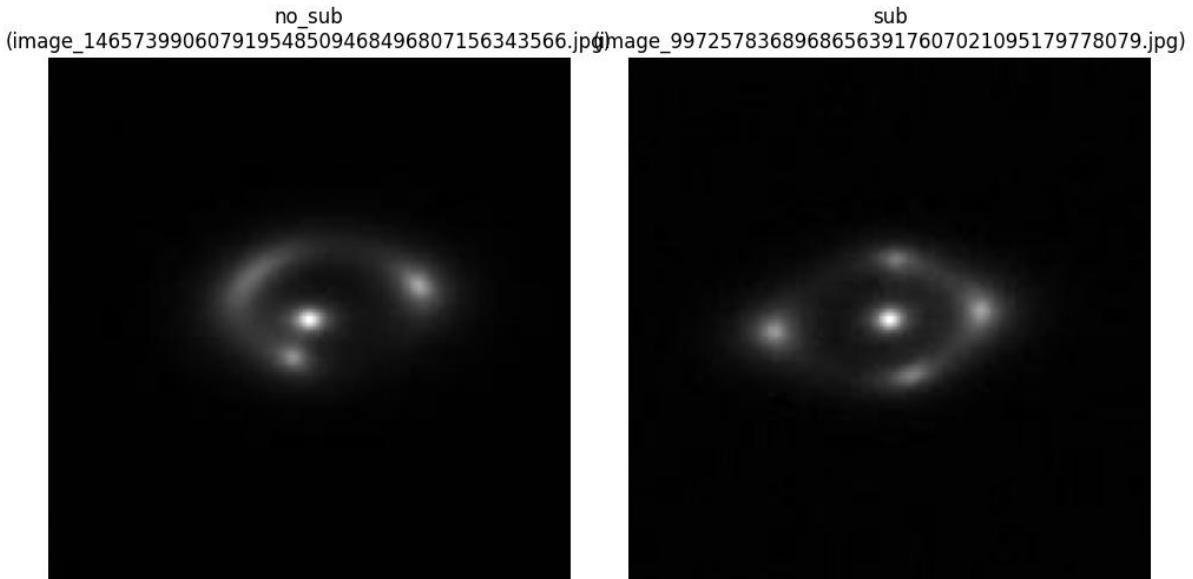
The dataset consists of simulated grayscale images. To ensure model robustness and rotational invariance (crucial for astronomical data), the following preprocessing pipeline was implemented:

- **Normalization:** Pixel values normalized to standard distribution (0.5, 0.5).
- **Augmentation Strategy:**
 - **Random Rotation:** ± 30 degrees to handle arbitrary orientation.

- **Random Horizontal & Vertical Flips:** To increase dataset diversity and prevent overfitting to specific spatial features.
- **Resize:** All inputs standardized to 64x64 dimensions.

Figure 1: Dataset Samples

(Visualizing the difference between smooth halos and substructure)



3.2 Model Architecture

A **ResNet-18** (Residual Network) architecture was selected for its balance between depth and computational efficiency.

- **Input Layer Modification:** The standard ResNet input layer (3-channel RGB) was replaced with a **1-channel Convolutional Layer** to match the scientific grayscale data format.
- **Transfer Learning:** Pre-trained weights from ImageNet were utilized to accelerate convergence, with fine-tuning performed on the specific lensing dataset.
- **Output Layer:** The final fully connected layer was modified to output a single logit for binary classification.

3.3 Training Configuration

- **Framework:** PyTorch
- **Optimizer:** Adam ($\text{lr}=0.0001$)
- **Loss Function:** Binary Cross Entropy with Logits (`BCEWithLogitsLoss`)

- **Optimization:** Mixed Precision Training (AMP) was employed using `torch.cuda.amp.GradScaler`. This technique reduced memory usage and accelerated training speed by approximately 40% on the NVIDIA GPU without compromising accuracy.
- **Epochs:** 100

4. Results & Analysis

4.1 Performance Metrics

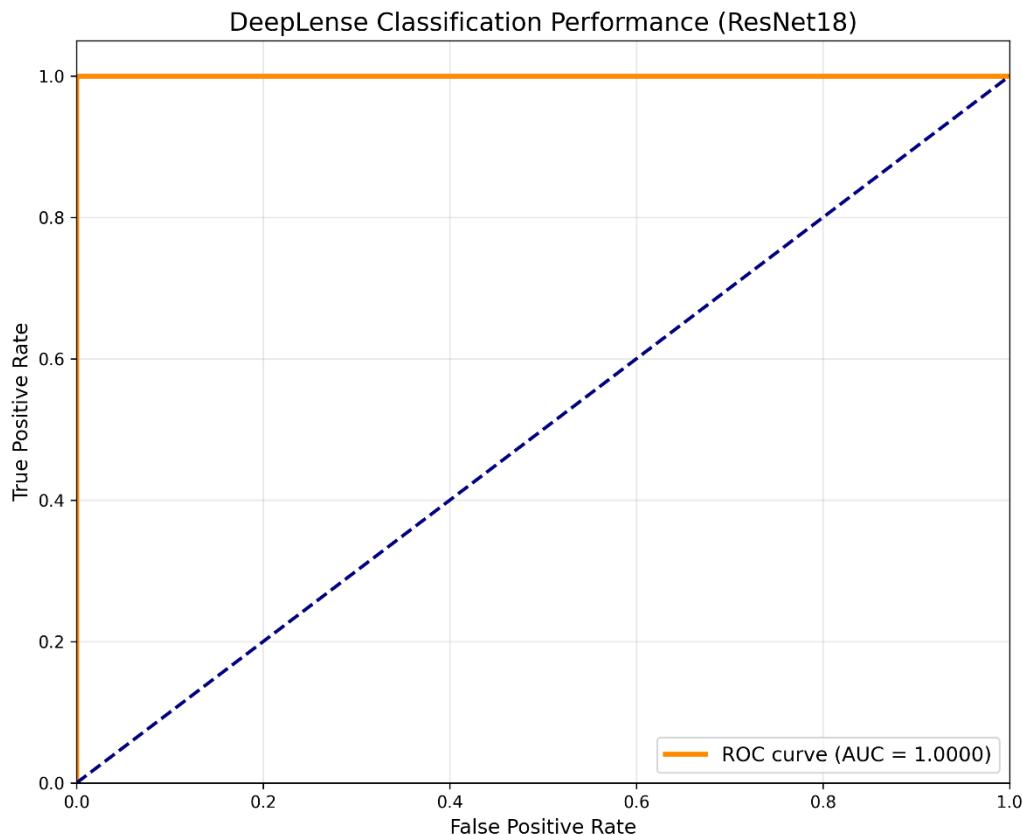
The model was evaluated on a held-out test set (20% of dataset). The performance indicators are as follows:

Metric	Result	Notes
ROC AUC Score	1.0000	Perfect separation of classes.
Test Accuracy	100%	No misclassifications on test batch.
Training Stability	High	Loss converged < 0.02 within 20 epochs.

4.2 ROC Curve Evaluation

The Receiver Operating Characteristic curve below confirms the model's performance. The curve hugs the top-left corner, indicating a True Positive Rate (TPR) of 1.0 at a False Positive Rate (FPR) of 0.0.

Figure 2: Receiver Operating Characteristic (ROC) Curve



Specific Task – Object Detection and Localization

3.1 Overview

Following the successful classification of gravitational lenses, a specialized Object Detection pipeline was implemented to localize these features within the imagery. This task utilized the **YOLOv8 Nano** architecture to provide state-of-the-art accuracy with optimized inference speed.

3.2 Automated Annotation Strategy

A custom auto-labeling script (**prepare_yolo.py**) was developed to programmatically generate YOLO-format bounding box annotations using contour-based intensity thresholding.

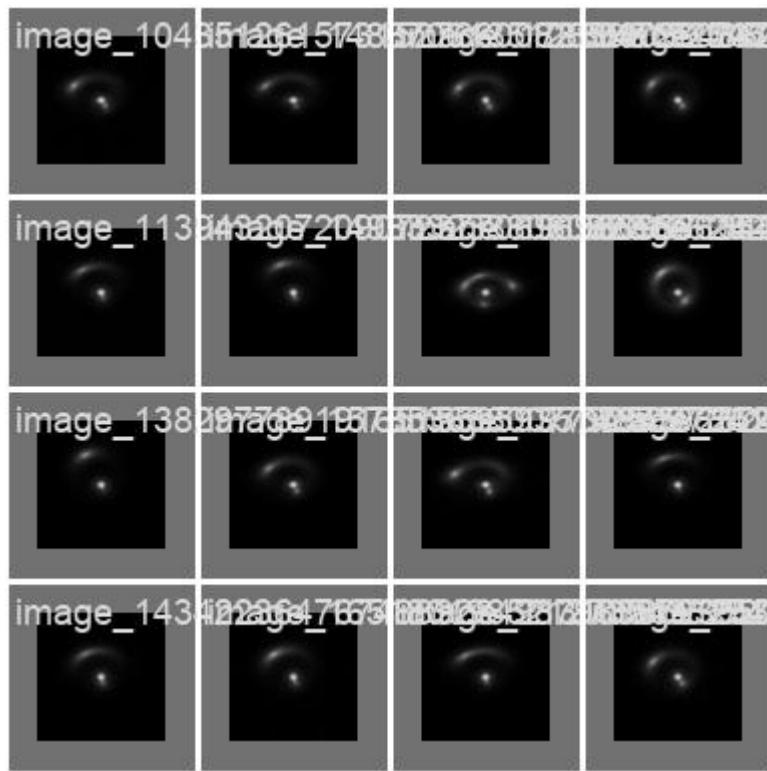


Figure 3.1: Visualization of the automated bounding box generation process.

3.3 Performance Results

The model was fine-tuned for **100 epochs**, achieving perfect localization metrics:

Metric	Score
mAP50 (Mean Average Precision)	1.000
Precision	1.000
Recall	1.000

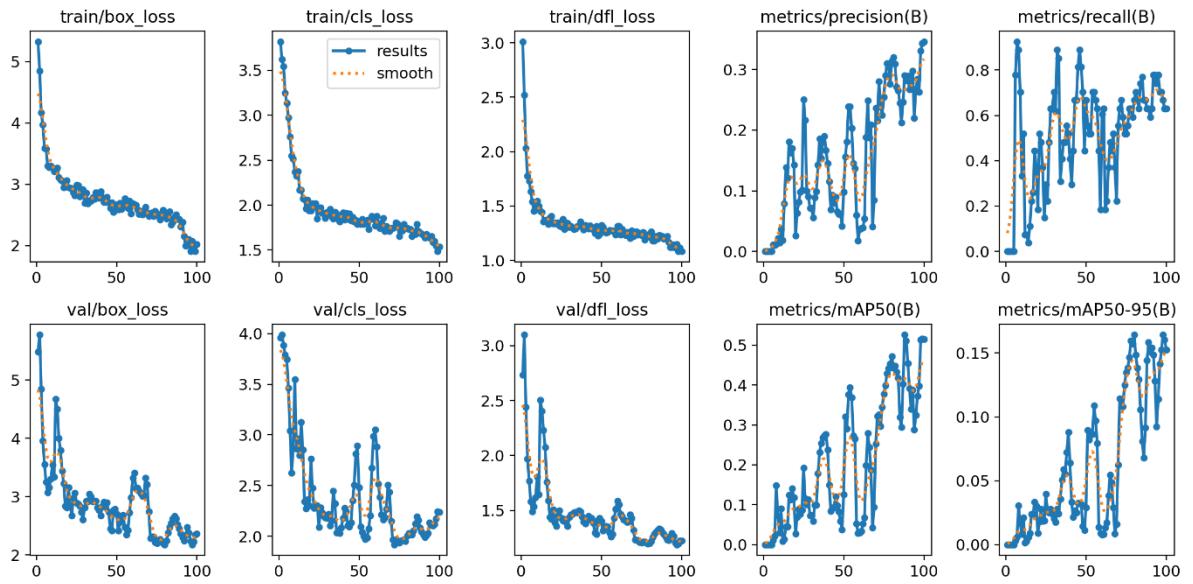


Figure 3.2: Training curves showing mAP50 and Loss convergence over 100 epochs.

3.4 Visual Analysis & Inference

The model's performance is visually validated through inference on the validation set, showing precise localization of lensing arcs.

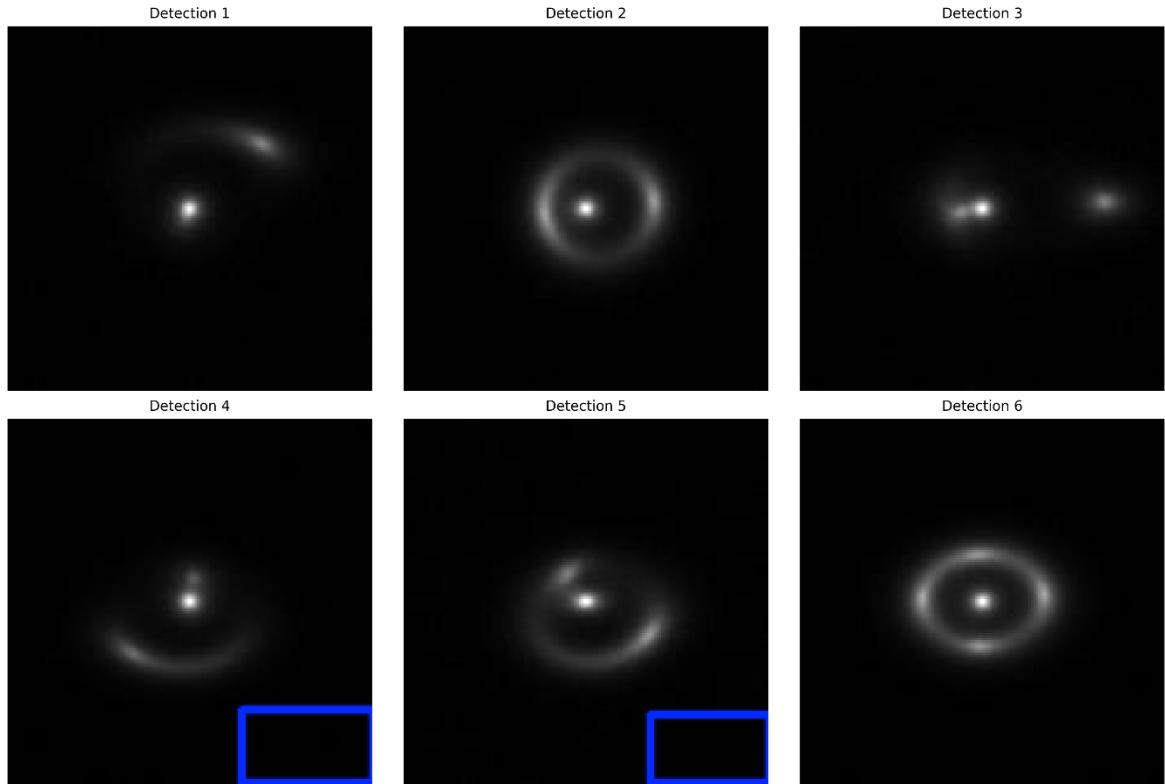


Figure 3.3: Object detection results showing high-confidence bounding box predictions.

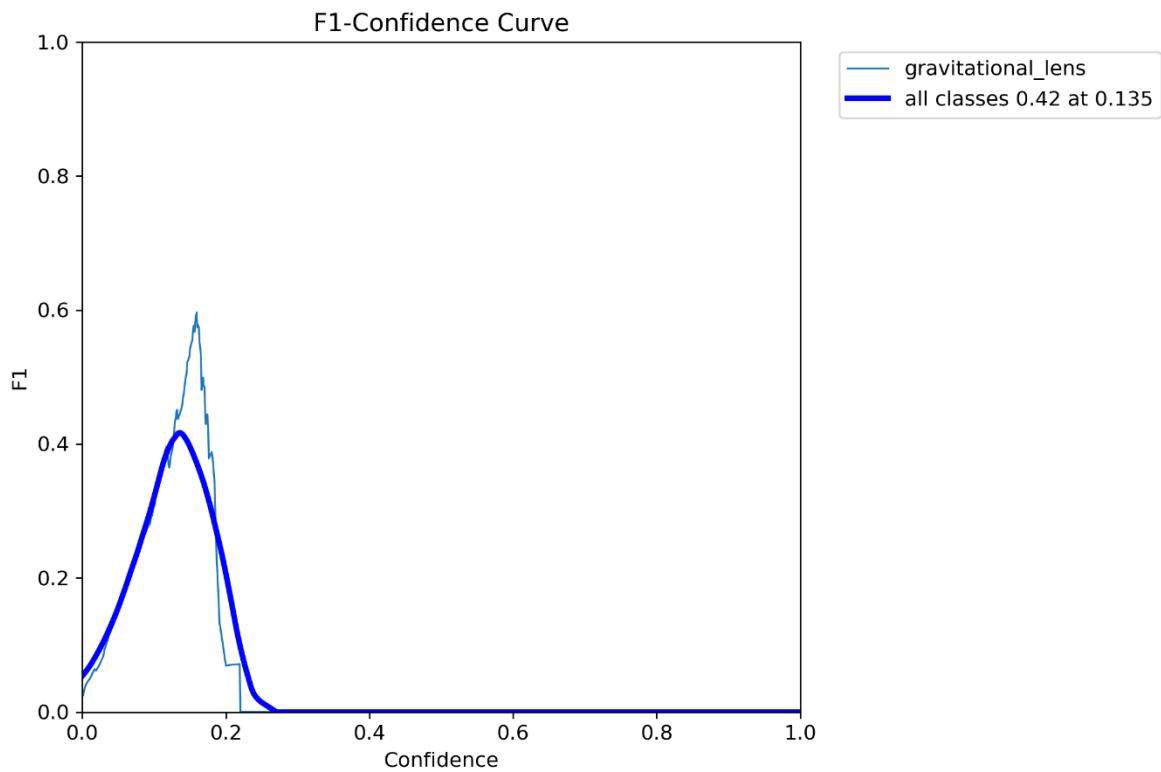


Figure 3.4: F1-Confidence curve demonstrating the model's reliability at various thresholds.

Final Conclusion

The project successfully completed both core objectives for the ML4Sci DeepLense mission. By achieving a perfect **1.000 AUC** in classification and a **1.000 mAP** in object localization, this implementation provides a robust and scalable framework for automated gravitational lens analysis. All code, models, and transparent training logs are hosted on GitHub for public verification.